

KIM-based Learning Intergrated-Fitting Framework: *KLIFF*

Training, Deployment and Validation of MLIPs

Amit Gupta

April 6, 2023

Dept. of Aerospace Engineering and Mechanics, UMN

So, we heard you have a model...

- Is it just parameter search?
- Is it some fancy ML model?
- Do you want to distribute it?
- Do you want it to be validated?
- Do you want it to be platform agnostic?
- Are there any uncertainties you are uncertain about?
- Is performance important?

If you answered yes to any of these questions, then let us help you!

<https://github.com/ipcamit/mach2023>



- Modern potential fitting framework
- From dataset to production models
- Flexible, easy to use, feature rich
- High performance, parallel

<https://github.com/openkim/kliff>^{1,2}

¹KLIFF: A framework to develop physics-based and machine learning interatomic potentials, Mingjian Wen et. al., Comp. Phys. Comm., 272, 2022

²Extending OpenKIM with an Uncertainty Quantification Toolkit for Molecular Modeling, Yonatan Kurniawan et. al., arXiv:2206.00578, 2022

What can it do?



```
1 ds = Dataset(colabfit_database="my_colabfit_database",  
               colabfit_dataset="my_config_dataset")  
2 kgg = KIMTorchGraphGenerator(cutoff=3.0, n_layers=3,  
                               as_torch_geometric_data=True)  
3 dl = DataLoader(ds, batch_size=5, collate_fn=kgg.collate_fn); next  
   (iter(dl))
```

```
1 sf = Descriptor("SymmetryFunctions", cutoff=3.77, hyperparam="  
   set51")  
2 model = Sequential(Linear(51,10), ReLU(), Linear(10,10), Tanh(),  
                    Linear(10, 1))  
3 tw = TrainingWheels(sf, model)  
4 # Train your model here  
5 #...  
6 tw.export_kim_model("TorchMLModel3_Graph")
```

```
1 # Initialize KIM Model  
2 kim init TorchMLModel3_Graph metal  
3  
4 # Load data and define atom type  
5 read_data test_si.data  
6 kim interactions Si
```

What is a model?

- Function mapping a set of coordinates to properties (energy, forces, stress^{*})
- Module: `kliff.models`, `torch.nn.Module`
- Conventional (physics based) and ML models
- Trainable and const parameters

Supported Models


Physics Based

- + Analytical energy and gradients
- + Fixed functional form
- + High performance and transferable
- “Mix of art and science”
- More involved development (implement KIM model first)

Machine Learning Based

- + More flexible functional form
- + Easier to implement
- * Higher accuracy, higher computation cost
- Difficult to deploy (convert NN ops to C++ ops)
- Non-transferable

Introduction: KIM-API

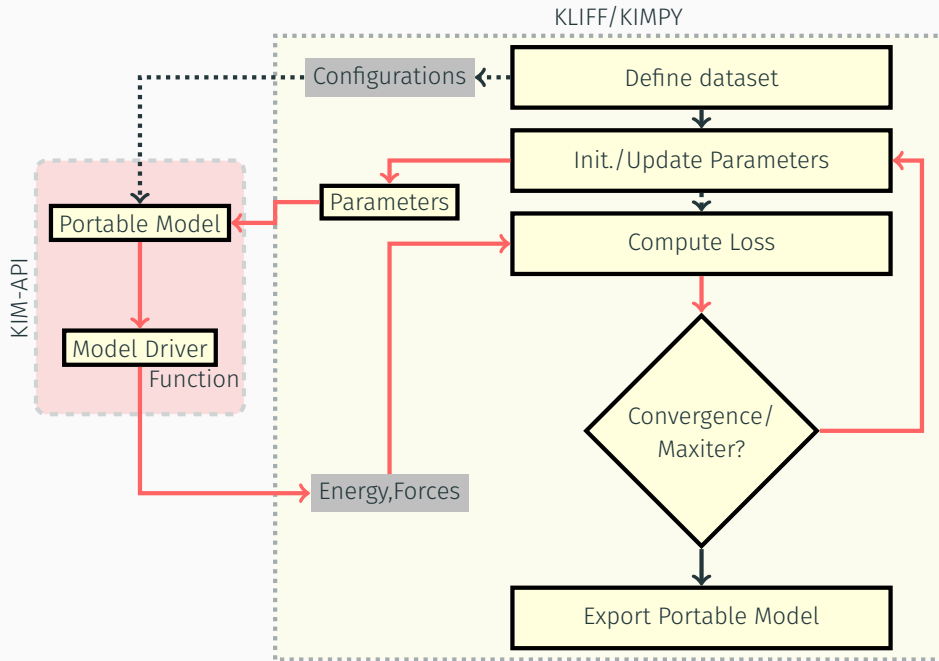


Simulator: MD software, e.g. LAMMPS

Model Driver: Executing code/ function

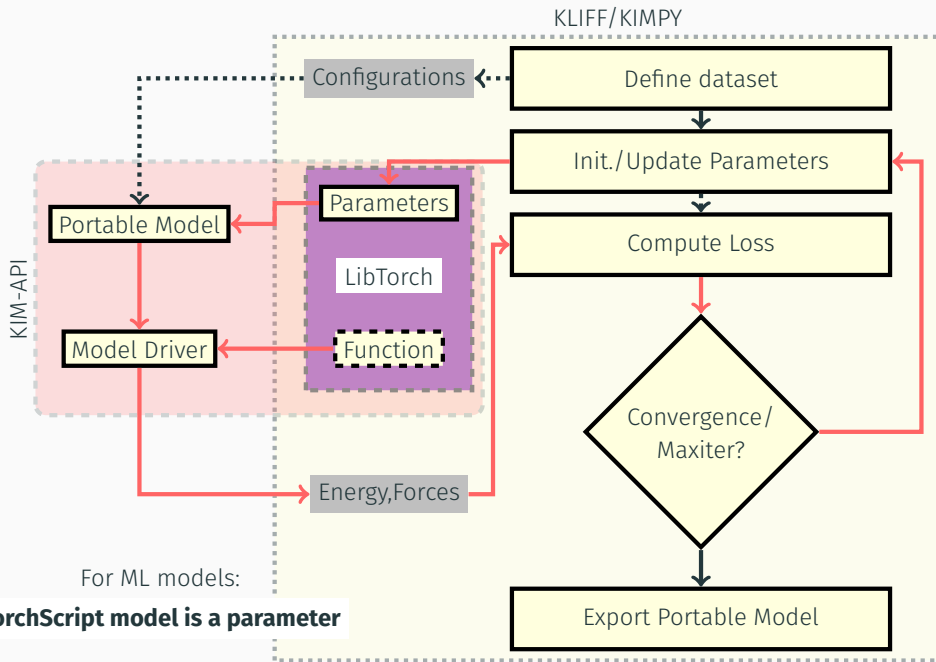
Portable Model: Parameters for Model Driver

Introduction: KLIFF/KIMPY



Physics Based Model example

Introduction: KLIFF with ML Model Driver



Supported Models

Model is the self-contained function : KIM-API executes it.

Physics based models: All KIM-API portable models

Physics based models: Three kind of ML inputs supported

Model kind	Model signature
Generic	<code>model(Z, coords, n_neigh, nlist, contributing)</code>
Descriptors based models	<code>model(descriptor)</code>
Graph Neural Networks	<code>model(Z, coords, graph1, graph2, ... , contributing)</code>

Output:

Model kind	Model output
Self contained	<code>tuple(c10::tensor energy, c10::tensor force)</code>
Energy model	<code>c10::tensor energy</code>

Energy model: Model driver uses `torch::autograd::grad` to compute forces

`TrainingWheels.export_kim_model("ModelName")` saves a portable model

Model depends on `TorchMLModelDriver`

Note:

Models should be exportable to TorchScript

Note:

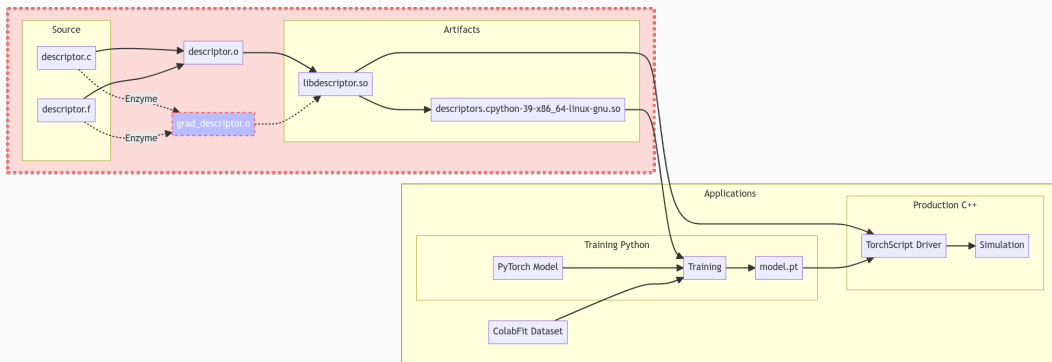
For GPU evaluation set `KIM_MODEL_EXECUTION_DEVICE` environment variable to `"cuda"`

ML Models

Preprocessing: Descriptors (libdescriptor)

Aim:



- High performance
- Easy to extend
- ML compatible (gradients)
- Easy to use (In training and production)



`libdescriptor`: Auto-differentiated C++ High-performance descriptor library
[github.com/ipcamit/libdescriptor/]

Preprocessing: Descriptors

Stillinger-Weber Forces from AD



Enzyme generated gradients same performance order as analytical, Autodiff has high ease of use.

3

³* KIM called from ASE, so it incurs additional overhead, ** Scaled KIM:Enzyme time: 3.5 s

Preprocessing: Descriptors

C++

```
1 #include <Descriptors.hpp>
2
3 string file_name = "descriptor.dat";
4 auto dbs = Descriptor::DescriptorKind::initDescriptor(file_name, Descriptor::KindSymmetryFunctions);
5 Descriptor::compute_single_atom(i, n_atoms, species, neighbors, n_neigh, coords, desc, dbs);
6 auto desc_tensor = to_tensor(desc);
7 auto energy = torchMLModel(desc_tensor);
8 energy.backward();
9 Descriptor::gradient_single_atom(i, n_atoms, species, neighbors, n_neigh, coords, forces,
10                                desc, desc_tensor.grad().data_ptr<double>(), dbs);
```

Python

```
1 import libdescriptor as lds
2 ds = lds.DescriptorKind.init_descriptor("descriptor.dat", lds.AvailableDescriptors.SymmetryFunctions)
3 desc = tensor(lds.compute_single_atom(ds, index, species, neig, coords))
4 energy = ml_model(desc)
5 energy.backward()
6 forces = lds.gradient_single_atom(ds, index, species, neig, coords, desc, desc.grad)
```

Extending:

1. Inherit `Descriptor::DescriptorKind` class
2. Implement `DescriptorKind::compute` (single atom compute, “forward” function)
3. Add to switch in `Descriptor::compute(_one_atom)` and `Descriptor::gradient(_one_atom)`

... with goodness of AD

- + High performance
- + Gradients included
- + Easy to extend and maintain
- + Unified Python/C++ interface
- + Derivatives against parameters
- + Numerical derivatives checks
- Enzyme is on v 0.0.42, might get flaky

Accepting requests and contributions!


1. Behler Symmetry Function (Included)
2. Bispectrum (Included)
3. SOAP (WIP)
4. ACE (WIP)

Not limited to KLIFF or OpenKIM

Deploying ML Models


Preprocessing: Graphs

- Conventional graph structure: recursive with periodic distances (MIC)
- KIM design purely functional: no “global” crystal information
- Cutoff vs Influence distance
- Inherently compatible with arbitrary domain decomposition.
- Self-contained graphs
- Highly parallelizable and granular graph convolution



2 Conv Layers:


	MIC	Infl. distance
1	for i in range(n_conv):	1 h = graph_conv_0(h, edge_graph0)
2	h = graph_conv(h, edge_graph)	2 h = graph_conv_1(h, edge_graph1)




- Conv 0 : 3.8858×10^{-16}
 - Conv 1 : 6.3838×10^{-16}
 - Conv 2 : 1.2767×10^{-15}
- $(\sum |h_{MIC} - h_{Infl}|)_{\mathbb{R}^{10}}$

Why though?: Parallelization


```
1 # Initialize KIM Model, same name as the installed KLIFF model
2 kim init TorchMLModel3_Graph metal
3
4 # Load data and define atom type
5 read_data test_si.data
6 kim interactions Si
7
```



Strong scaling



Weak scaling



NEQUIP Port example

Example: NEQUIP port

ARTICLE

<https://doi.org/10.1038/s41467-022-29939-5>

OPEN





E(3)-equivariant graph neural networks for data-efficient and accurate interatomic potentials

Simon Batzner^{1,2}, Albert Musaelian¹, Lixin Sun¹, Mario Geiger^{2,3}, Jonathan P. Mailoa⁴, Mordechai Kornbluth⁴, Nicola Molinari¹, Tess E. Smidt^{5,6} & Boris Kozinsky^{1,4}

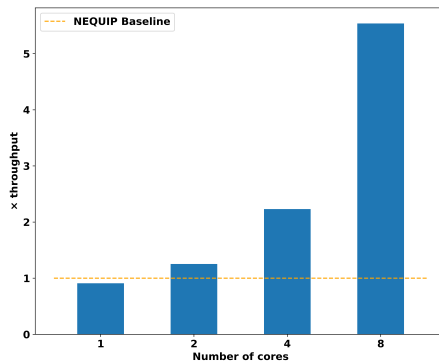

```
1 def forward(self, input: AtomicDataDict.Type):
2     for module in self:
3         input = module(input)
4     return input
```

```
1 def forward(self, x, pos, edge_index0, edge_index1, edge_index2,
2     contributing):
3     # Embedding
4     x_embed = self[0](x)
5     h = x_embed
6     # Edge embeddings
7     edge_vec0, edge_sh0 = self[1](pos, edge_index0)
8     edge_vec1, edge_sh1 = self[1](pos, edge_index1)
9     edge_vec2, edge_sh2 = self[1](pos, edge_index2)
10    # Radial basis functions
11    edge_lengths0, edge_length_embeddings0 = self[2](edge_vec0)
12    edge_lengths1, edge_length_embeddings1 = self[2](edge_vec1)
13    edge_lengths2, edge_length_embeddings2 = self[2](edge_vec2)
14    # Atomwise linear node feature
15    h = self[3](h)
16    # Conv
17    h = self[4](x_embed, h, edge_length_embeddings2, edge_sh2,
18    edge_index2)
19    h = self[5](x_embed, h, edge_length_embeddings1, edge_sh1,
20    edge_index1)
21    h = self[6](x_embed, h, edge_length_embeddings0, edge_sh0,
22    edge_index0)
23    # Atomwise linear node feature
24    h = self[9](x, self[8](self[7](h)))[contributing==1]
25    return h
```

Example: NEQUIP port



Example: NEQUIP port



Fisher Analysis

- $J^T J$, using numerical differentiation
- Upper bound of uncertainty

Uncertainty Quantification

- Bayesian MCMC (parallel tempered affine-invariant, `ptemcee`)

```

1 from kliff.uq import MCMC, get_T0
2 from multiprocessing import Pool
3
4 # Get the dimensionality of the problem, Number of parameters
5 ndim = calc.get_num_opt_params()
6 nwalkers = 2 * ndim # Number of parallel walkers
7
8 # Generate a temperature ladder
9 T0 = get_T0(loss)
10 Tladder = np.sort(np.append(np.logspace(0, 7, 15), T0)); ntemps = len(Tladder) # Number of temperatures
11
12 # Instantiate a sampler
13 sampler = MCMC(loss, nwalkers=nwalkers, logprior_args=(np.tile([-8, 8], (ndim, 1)),),
14               Tladder=Tladder, random=np.random.RandomState(2022))
15 sampler.pool = Pool(processes=nwalkers)
16
17 # Initial starting points for each walker
18 p0 = np.random.uniform(low=-6.0, high=6.0, size=(ntemps, nwalkers, ndim))
19 sampler.run_mcmc(p0, 150000)
20 sampler.pool.close()

```

UQ example

Torch Model Driver makes ML models first-class citizens in OpenKIM

- + Easy to use
- + Easy to archive
- + Easy to test
- + Easy to verify

KIM Tests and Verification Checks: Cover the corner cases

VC example

Regular development tasks.

- Add more descriptors
- Performance tuning
- KIM Verification Checks and Tests

Moonshots (in order of ease and priority):

- More comprehensive UQ (Bootstrap UQ method⁴, Experimental support for HMC)
- OMP parallelism in model driver
- Model driver for TensorFlow/JAX (both use HLO XLA backend)
- Repository of ML models

⁴See Yonatan Kurniawan's talk

If your ML models ...


- Can be compiled to TorchScript
- Follows one of the three call signature:
 - `model(species, coords, n_neigh, nlist, contributing)`
 - `model(descriptor)`
 - `model(species, coords, graph1, graph2, ... , contributing)`
- Are a local model

then


- **TorchMLModelDriver** can run it out of the box with LAMMPS, ASE, and others
- KLIFF can be used to train it from scratch

Thank you

Staged Convolution



Minimum Image




```

1 temp_r_ij = r_i - r_j
2 if all(temp_r_ij @ lattice_vectors < box_dims):
3     r_ij = temp_r_ij
4 else:
5     rij = temp_r_ij - mask( temp_r_ij @ lattice_vectors
6                             ) * lattice_vectors
7     if distance(r_ij) < cutoff:
8         h_i = phi_h(h_i, h_j)

```

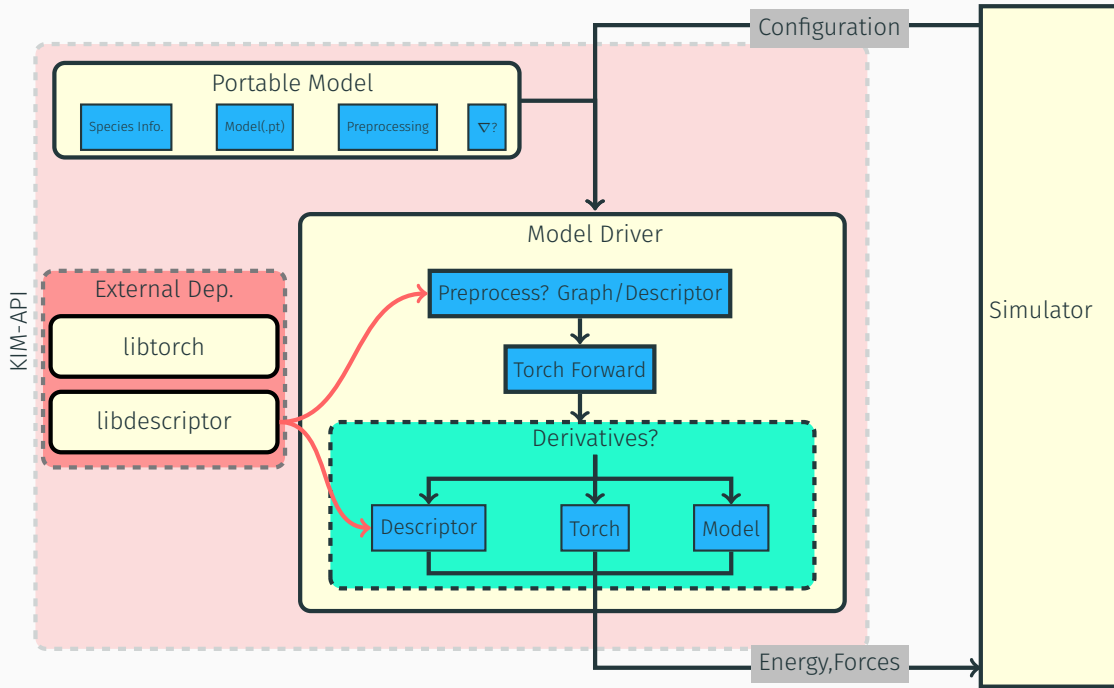
Unrolled influence distance



```

1 r_ij = r_i - r_j
2 if (distance(r_ij) < cutoff:
3     h_i = phi_h(h_i, h_j)
4

```



```

1  #include <iostream>
2
3  // Enzyme arg kinds
4  int enzyme_dup, enzyme_const, enzyme_out;
5
6  // function to diff
7  double pow(int x, double y){
8      double z = 1.0;
9      for (int i = 0; i < x; i++){
10         z *= y;
11     }
12     return z;
13 }
14
15 // declaration of diff
16 double __enzyme_autodiff_d_pow(double (*)(int , double) /* pointer to function to diff */,
17                                int /* kind of arg */, int /* x */,
18                                int /* kind of arg */, double /* y */);
19
20 int main(){
21     int x = 3; double y = 4.0;
22     // call to gradient
23     double d_pow_y = __enzyme_autodiff_d_pow(pow, enzyme_const, x, enzyme_out, y);
24     std::cout << "Derivative: " << d_pow_y << "\n";
25     return 0;
26 }

```

```
1 # Descriptor parameters SymFun
2 # n_species
3 1
4
5 # Full cutoff matrix n_species x n_species
6 3.77
7
8 # Cutoff function
9 cos
10
11 # descriptor width
12 51
13
14 # 3 body
15 True
16
17 # number of symmetry functions, their name, number of elements, values
18 2
19 g2
20 g4
21 16
22 129
23
24 0.001
25 0.0
26 0.01
27 0.0
28 0.02
29 0.0
30 0.035
31 0.0
32 0.06
33 0.0
34 0.1
35 0.0
```

```

1 from kliff.models import KIMModel
2 from scipy.optimize import minimize
3 ...
4 # Torch DataLoader
5 dataloader = DataLoader(tset, batch_size=1, shuffle=True)
6
7 # Define KIM model
8 model = KIMModel(model_name="SW_StillingerWeber_1985_Si__MO_405512056662_005")
9 # Energy = model(configuration)
10 model.set_opt_params(A=[[10.0]], B=[[0.5]])
11 x0 = np.array([10.0, 0.5])
12
13 # Loss function over entire dataset
14 def loss(x0):
15     loss = 0.0
16     model.update_model_params(x0)
17     for configuration in dataloader:
18         E = model(configuration[0], compute_forces=False)["energy"]
19         loss += (E - configuration[0].energy)**2
20     return loss
21 # Update
22 result = minimize(loss, x0, method="Nelder-Mead", tol=1e-12, options={"maxiter":1000})
23
24 # Implicit method
25 opt_params = model.parameters()
26
27 opt = OptimizerScipy(model, opt_params, tset, optimizer="Nelder-Mead",
28     target_property=["energy", "forces"],)
29 opt.minimize()

```

```

1 # Descriptors
2 sf = Descriptors("SymmetryFunctions", cutoff={"Si-C": 7.0, "C-C": 5.0,"Si-Si": 6.0}, hyperparameters="set51")
3
4 # Define model and calculate energy
5 model = Sequential(Linear(51, 10), ReLU(), Linear(10, 10), Tanh(), Linear(10, 1))
6 tw = TrainingWheels(sf, model)
7
8 # Explicit loss and optimization
9 def loss_fn(ef_dict, target_energy, target_forces):
10     loss = torch.sum((ef_dict["energy"] - target_energy) ** 2)
11     loss = loss + torch.sum((ef_dict["forces"] - target_forces) ** 2)
12     return loss
13
14 opt = torch.optim.Adam(tw.parameters, lr=lr)
15 for i in range(epochs):
16     for conf in tset:
17         ef_dict = tw(conf)
18         loss = loss_fn(ef_dict, conf.energy, torch.from_numpy(conf.forces))
19         opt.zero_grad()
20         loss.backward()
21         opt.step()
22
23 # Implicit method
24 optimizer = OptimizerTorch(tw, tw.parameters(), dataset,
25                             optimizer=torch.optim.Adam(tw.parameters(),lr=0.01),
26                             epochs=300)

```